

DAY 16 / 30

THE FORGE · PRACTICAL GENAI

Retrieval-Augmented Generation (RAG) explained

Retrieval-Augmented Generation (RAG) explained

Remember Day 9: models hallucinate and their knowledge is frozen at a training cutoff. RAG is the most common fix, and it's more straightforward than the acronym suggests.

- ◆ Your question gets converted into an embedding (Day 8).
- ◆ The system searches a knowledge base — your company docs, a product manual, today's news — for chunks of text with the closest matching embeddings.
- ◆ Those retrieved chunks get inserted into the prompt as context.
- ◆ The model generates its answer grounded in that retrieved text, instead of purely from memory.

DAY 16 / 30

“Any tool you use daily is almost certainly RAG under the hood — customer support bots, internal doc search, "ask your codebase" tools.”

Follow along — Day 17 drops tomorrow.

#RAG

#AI

#LLM

#GenerativeAI