

DAY 26 / 30

SUMMIT RIDGE · ML ENGINEERING

Fine-tuning vs prompting: when to customize a model

Fine-tuning vs prompting: when to customize a model

You want a model to behave a specific way for your use case. Do you write a better prompt, or retrain the model itself? This is one of the most common strategic decisions teams building on AI face — and prompting wins far more often than people assume.

- ◆ Prompting (including RAG from Day 16): you shape behavior entirely through the input, without touching the model's weights. Fast to iterate, cheap, easy to update, and surprisingly powerful for most tasks — including teaching a model new "knowledge" via retrieval.
- ◆ Fine-tuning: you continue training an existing model on your own examples, actually adjusting its weights (Day 3) so the new behavior becomes baked in. This is slower, more expensive, and requires real ML expertise — but it's the right call when you need a very specific output format at scale, a distinct tone/style baked in reliably, or behavior that's hard to reliably coax out through prompting alone.

“Skipping straight to fine-tuning when a better prompt would've solved it is one of the most common wastes of time and budget I see on AI projects.”

Follow along – Day 27 drops tomorrow.

#MachineLearning

#AI

#FineTuning

#GenerativeAI