

DAY 29 / 30

SUMMIT RIDGE · ML ENGINEERING

Cost & compute: the economics of running AI

Cost & compute: the economics of running AI

The most surprising thing to newcomers building their first AI product isn't a technical concept — it's the bill. Understanding where AI costs actually come from changes how you design a product from day one.

- ◆ Training cost — a one-time (or periodic) expense, often enormous for large models, but amortized across every future use.
- ◆ Inference cost — what you pay every single time the model runs (Day 7). For a product with real usage, this dwarfs training cost over time — it's the ongoing "cost of goods sold" for an AI feature, not a one-off R&D expense.
- ◆ The levers that actually move the number: model size (bigger models cost more per call), output length (you're often billed per token generated), prompt length (more context in = more cost), and call frequency (does every user action trigger a model call, or only the ones that need one?).

“The teams that treat inference cost as a design constraint from the start build meaningfully more sustainable AI products than the ones who bolt on efficiency after launch.”

Follow along – Day 30 drops tomorrow.

#AI

#MLOps

#Cost

#GenerativeAI